

# Themis.ai

The matter operating layer between general AI, legal authority, Word, and the decision that actually gets recorded.

203


Reddit threads of primary buyer research

665

backend tests passing on a working system

5

model providers — the lawyer keeps theirs



Today — the ranked briefing the lawyer opens on

# Legal work starts in ambiguity. The setup is the tax.

A request reaches product counsel as a Slack thread, a spec, a screenshot, and a deadline. Before any judgment is applied, the lawyer performs the same reconstruction — every time, by hand.

*Every one of these happens before the lawyer does the one thing only the lawyer can do.*

1

Work out what the business is actually trying to do

2

Separate the material facts from the missing ones

3

Find the prior decision and playbook that already govern this

4

Decide what research is even necessary

5

Draft the response, then track what was decided

6

Revisit it when the law, the product, or the facts move

Existing tools store documents, manage tasks, or answer isolated questions. None of them organises the path from intake to decision the way a product lawyer actually thinks about the work.

# The buyers are already saying what is missing

- “We had access to each system for ~3 weeks... The only one that caught it without further prompting was Claude Opus.”

~30-attorney firm, after trialing CoCounsel, Harvey and Legora — then building its own tool instead. r/legaltech, Aug 30 2026

- “For 5 Harvey seats (minimum), 12 months lock-in, I was told it was USD \$425/seat/month.”

A 3-attorney legal department separately rejected roughly \$30,000/year of the same product. r/legaltech, Jul 20 2026

- “Each member just creates a Claude chat and uploads the contract, the playbook, and any additional context needed.”

The context tax, paid again on every single matter. r/legaltech, Apr 22 2026

- “I missed that AI had hallucinated an entire statutory provision.”

A reviewer cite-checked an AI memo and still missed it. Proof is the product, not prose. r/biglaw, Jun 6 2026

# Legal AI is not one market. Buyers assemble a stack.

BUYER JOB

WHO WINS IT TODAY

WHY

**Think, draft, explain**

ChatGPT / Claude

General reasoning, fast writing, low entry cost

**Find legal authority**

Westlaw / Lexis

Known corpus, links to authority, the citator habit

**Repeatable document work**

Harvey / Legora / CoCounsel

Review tables, document scale, Word and DMS

**Apply a firm's own method**

An internal build

Exact playbooks, precedent, connectors, control

**Remember and move the matter**

**No clear owner for small teams**

Practice systems store records. None create a foothold or a next action.

*Most buying disappointment comes from expecting one product to do all five jobs. The fifth is the opening.*

# One matter, from a pasted request to a recorded decision

- 1

INTAKE ONCE

Paste a Slack thread, an email or a redline note. Drop the files. Keep your own model.
- 2

ORIENT AUTOMATICALLY

Business objective, material facts, missing facts, issues, the prior decisions that already govern it, and one next action.
- 3

DO THE FIRST WORK


Targeted research or document review, queued per matter, saved into the matter as an editable record.
- 4

PROVE IT

Agent output stays visibly agent output — labelled, dashed, source status shown — until a human accepts it. Passage-level proof is the next build.
- 5

HAND IT OFF

A review-ready draft, a recorded decision with a named decider and a revisit date, and a next action that survives the model.



The lawyer never reconstructs the matter to ask the next question.

# Everything needed to act, on one page

The screenshot shows the themis.ai web application. The top navigation bar includes the logo, a user profile 'BH', and links for Today, Briefing, Workspace, Matters (active), Decisions, Skills, Automations, Agents, and Settings. Below the navigation bar, the main header displays 'Overdue' and 'Matters' tabs, followed by the matter title 'Relay: Open-Banking Reauthentication and Data Use'. On the right side of the header, there are filters for STAGE (Waiting on your judgment), RISK (Set risk), and DUE (0 days late). The left sidebar shows 'Matter contents' with a link to 'Matter overview'. The main content area is titled 'Overview' and contains a 'REQUIRED WORK' section. This section includes a heading 'Select the customer reauthentication path', a stage description 'Stage: Waiting on your judgment. Brian Harris must act next: Select the customer reauthentication path', a 'Matter at a glance' box with background information, a 'Question to resolve' box with a specific query, an 'Original request' box, and a 'Participants' list with roles: Omar Haddad (Requester), Brian Harris (Legal owner), and Nina Foster (Business owner). There is also a 'Participant name' input field at the bottom.

## The question, in plain language

Not a file tree. The business objective and the exact question to resolve sit above everything else.

## Who must act next, by name

Ownership and lateness are first-class. “Waiting on Brian Harris” is a state the system computes, not a status someone remembered to set.

## Judgment and work product are separate acts

Record decision and Finalize draft are two explicit buttons. Neither happens because an agent wrote something.

## The customer owns the records

Every matter is plain Markdown on disk. SQLite is a disposable index that rebuilds. There is no database to be locked out of.

# The second lawyer in the room

In-house counsel is structurally alone. There is no partner down the hall, no one to say “have you thought about this?” — and every blind spot is theirs alone. Themis is not there to hand down the answer. It is there to make sure the lawyer sees the whole question before they give theirs.

01

## Sees the space with you

Maps the business objective, the live issues, the realistic paths, and the few facts that actually move the outcome. It widens what the lawyer has considered rather than replacing their consideration.

02

## Says where it stops

Names what it assumed, what it could not determine, and what it never read. A colleague who admits the edge of their knowledge is worth more than one who is never unsure — and lawyers already say so.

03

## Remembers why

Holds the conversation, the reasoning, the recorded decision and its date. A year later it can say what was decided, on what basis, and what has changed underneath it since.

## WHY WE BELIEVE THIS

■ **30% of 203 threads** are lawyers asking other lawyers what is normal — improvising the colleague they do not have.

■ **“...it asks for clarification when it doesn’t know.”** The best-rated tool in the corpus wins on colleague behaviour, not on answer quality.

*The lawyer still gives the answer. Themis makes sure it is given with the whole question in view — and remembers why, when somebody asks a year from now.*

# Four architectural bets, each one a reason not to churn

01

## Model-neutral by construction

Five provider adapters ship today, with per-agent model and reasoning-effort selection. The buyer keeps Claude, ChatGPT or an approved internal endpoint. Model quality moves too fast to be anyone's moat — so we refuse to make it ours.

02

## Markdown is the source of truth

Every matter, decision, agent, tool, workflow and schedule is a readable file on disk. SQLite is a rebuildable index. There is no lock-in to defend, which is exactly why a lawyer will trust it with the record.

03

## Recommendation and decision are separate objects

Competitors keep this distinction in a chat transcript. Themis keeps it in the schema, in the type system and on the screen. It is the difference between a tool and a system of record.

04

## The intelligence is editable text

Agents, tools, company context and playbooks are Markdown the lawyer can read and change — configurable enough to beat an internal build, and far cheaper to maintain than one.

# Highest setup tax, lowest rollout cost

## BEACHHEAD

**Solo GC and 1–5 person product / legal teams**

Highest fit

Constant ambiguous intake, no junior support, company-specific history to carry, and no appetite for an enterprise rollout. Can test on one real matter without migrating anything.

## ADJACENT

**5–30 lawyer boutiques and specialty firms**

High, if the workflow is narrow

Need repeatable document work and matter continuity, but reject seat minimums and broad suites.

## LATER

**30–100 lawyers with a technology lead**

Conditional

Must beat an internal build on configuration, DMS handoff and maintenance — the segment that most often builds its own.

## NOT NOW

**150+ lawyer enterprise and BigLaw**

Low for this product

Purchases turn on SSO, permissions, DMS, governance and support. None of that is built, and none of it is needed to prove the wedge.

Not a replacement for Westlaw or Lexis. Themis brings lawyer-checked authority into the matter; it never claims a proprietary corpus or a citator.

# A working system, deliberately narrow

## BUILT AND PASSING

- Next.js + FastAPI application, 16 product surfaces
- 665 backend tests green; typecheck and production build clean
- Markdown vault as source of truth, rebuildable SQLite index
- Five model providers with per-agent selection
- Research queues, decision register, staleness checks, scheduler
- Document review with native DOCX and PDF export

## DELIBERATELY NOT BUILT

- Auth, SSO, tenancy, team permissions
- Source-layout-preserving Word round trip
- Citator integration and court-grade cite validation
- Embeddings, queues, cloud deployment
- Multi-agent voting and reviewer vetoes
- Each excluded on purpose. None is needed to prove the wedge.

## WHAT WE PROVE NEXT

- Exact source passages plus a document coverage receipt
- Automatic matter dossier and next action on return
- Structured multi-document review with exceptions
- Review-ready Word handoff that keeps the support trail
- Lawyer-checked Westlaw / Lexis authority attached to a proposition
- Design partners doing real product-counsel work

*Themis.ai runs today as a single-user local web application. It has no authentication, tenancy or enterprise authorization, and is not for privileged material in an untrusted environment. That is the honest boundary of the current build.*

# The assumptions this rests on

The product's own discipline, applied to its deck. Each of these is load-bearing. Each names what would prove it wrong.

THE ASSUMPTION

WHAT WOULD PROVE IT WRONG

**Lawyers want a colleague, not an oracle.**

Every valued use in the corpus is companion-shaped — issue spotting, first pass, “where to start.”  
None is “give me the answer.”

■ Design partners who skip the edge map and want only the conclusion.

**Matter memory becomes real switching cost.**

The compounding bet: the register is worth more every month and cannot be imported by a competitor.

■ Teams that never revisit a decision, or that let the register go stale because keeping it current costs more than it returns.

**The small-team wedge is reachable before enterprise controls.**

No SSO, tenancy, DMS or ethical walls are needed to prove value on one real matter.

■ Prospects blocking on confidentiality and approved deployment before they will try it. The corpus already shows this risk live.

**Model neutrality stays an advantage.**

Model quality moves too fast to be anyone's moat, so we decline to make it ours.

■ One model plus harness becoming decisively better at legal work — at which point neutrality is a liability, not a virtue.

**The evidence base supports what we claim from it.**

203 de-duplicated, self-selected Reddit threads, March–August 2026, graded by confidence.

■ It establishes an unmet workflow pattern. It does not establish market size, willingness to pay, or that these posters are buyers.