

MARKET RESEARCH REPORT

What legal AI customers are actually saying

Customer voices on ChatGPT, Claude, Harvey, Legora, CoCounsel, and GC AI - how the tools are used, why they are chosen, and where the market still falls short

Prepared for	Legal AI market and product research
Date	September 1, 2026
Corpus	203 de-duplicated Reddit threads; 100 added sources from March-August 2026
Communities	r/legaltech (75), r/biglaw (39), r/Lawyertalk (30), r/LawFirm (23), and smaller lawyer-centered groups
Evidence standard	Direct use and task comparisons receive the most weight. Anonymous, vendor-linked, hearsay, and incomplete claims remain directional.

Bottom line

There is no single legal AI winner. Lawyers assemble a stack: ChatGPT or Claude for broad thinking and drafting; Westlaw, Lexis, or another controlled source for authority; and specialist tools for selected document, research, or team workflows. Products win when they remove work after the first answer. They lose when they add cost, verification, or another disconnected system.

The market in one page

The customer voices show an active but unsettled market. Lawyers are using AI every week, but use is concentrated in work that is easy to review: first drafts, summaries, issue lists, correspondence, clause work, document extraction, and research starting points. Trust falls as the work becomes citation-heavy, fact-sensitive, strategic, or difficult to verify.

Market pattern	What customers report	What it means
General models are the baseline	ChatGPT appears in 85 records and Claude in 73. Users value breadth, current model quality, flexibility, and price.	A specialist product must add workflow value, not only a legal label.
Specialists win selected jobs	Harvey, Legora, and CoCounsel are valued for legal sources, review tables, large files, Word/DMS links, and approved deployment.	Operational leverage is the main reason to pay a premium.
Research stays controlled	Users commonly use AI for orientation, then verify in Westlaw, Lexis, Google, law-firm memos, or supplied authorities.	No product in the corpus removes the lawyer's verification duty.
Firm size changes the choice	Small teams resist seat minimums and rollout burden; larger firms pay for standardization, security review, integration, and easier adoption.	The same product can be good value for BigLaw and poor value for a three-lawyer team.
Users build stacks	Examples combine Claude with Westlaw or Lexis, ChatGPT Enterprise with point tools, and specialist review with Word or DMS systems.	The practical competitor is often a bundle, not one vendor.

Discussion coverage

Product	Threads mentioning it	How to read the count
ChatGPT	85	Broad baseline across drafting, research, client behavior, policy, and operations.
Claude	73	Strong recent discussion around drafting, document work, projects, skills, and internal workflows.
Harvey	60	Frequent enterprise and legaltech comparisons; substantial positive and negative evidence.
CoCounsel	30	Mostly research, litigation, records, discovery, and Westlaw-related workflow.
Legora	29	Concentrated in BigLaw, review-table, drafting, diligence, research, and vendor-choice discussions.
GC AI	5	Evidence is too sparse for a strong market conclusion.

These are discussion counts, not usage, satisfaction, preference, or market-share measures.

How lawyers are using legal AI

Use is strongest where the output can be checked quickly against known facts, supplied documents, or an existing draft. The same users who report major time savings often reject AI as final authority.

Work category	Observed uses	Reported value	Recurring failure
Drafting and rewriting	First drafts, clauses, correspondence, advice, discovery requests, outlines, defined terms, cross-references.	Faster starting point and less typing; sometimes hours saved.	Generic language, wrong assumptions, weak strategy, or a second editing job.
Summaries and orientation	Cases, transcripts, contracts, records, bundles, call notes, business requests.	Rapid issue map and easier entry into unfamiliar material.	Missing material facts, overlong summaries, or false confidence that all files were read.
Legal research	Issue framing, source leads, case analysis, citation-linked starting points.	Useful signposting and faster first pass when sources are constrained.	Invented citations, wrong holdings, missing currentness, and mandatory manual checks.
Document-scale review	Diligence, discovery, timelines, testimony search, review tables, extraction, comparisons.	High value when the system handles many files and exposes results in a table.	Skipped concepts, incomplete coverage, timeouts, and unclear exception handling.
Firm-specific workflow	Playbooks, precedent cleanup, NDA review, prompt libraries, internal connectors, matter templates.	Better results after context and examples are supplied.	Setup burden, maintenance, rigid vendor workflows, or poor fit with existing systems.
Operational work	Intake, meeting notes, task preparation, pitches, client communications, simple automation.	Broad nonlegal value improves subscription economics.	Fragmented tools and weak handoff into the matter or practice system.

The common working stack

- General AI for thinking, drafting, summarizing, and mixed legal-commercial work.
- Westlaw, Lexis, PACER, supplied cases, or another controlled source for authority and citation checking.
- Word, a DMS, a CLM, or a point tool for the actual review and document workflow.
- Human review for legal judgment, material facts, citations, and final work product.

Customer behavior

The market is not moving from 'no AI' to 'one AI platform.' It is moving toward mixed stacks in which each tool is trusted for a different part of the work.

Why customers choose one tool over another

The choice changes with the work. General models win when drafting quality, flexibility, and price matter most. Legal products win when sources, Word or DMS workflow, review tables, and firm-wide adoption remove more work than the model alone. The examples below show the contrast in actual customer decisions.

Use case / customer	Chosen over	Why the customer chose it	Customer evidence
Reasoning, citations, Word documents, and redlines; 3-attorney legal department	ChatGPT over Harvey and Claude	ChatGPT performed best on the same prompts. Harvey's uneven issue coverage did not justify about \$30,000 per year.	Reported result: ChatGPT was preferred for every tested work category. (reddit-110)
Transactional drafting and precedent cleanup; lawyer and partner commenters	Claude over CoCounsel for generation	Claude was strong at defined terms, cross-references, precedent cleanup, and first drafts when supplied with the deal context.	"For the dirty work of drafting, it's insanely good." (reddit-131)
Corporate legal deployment; team size not stated	Harvey over ChatGPT	The buyer valued references, nuanced examples, Word integration, legal databases, and a usable interface more than a clear answer-quality lead.	"We couldn't prove that the AI answers are consistently better than ChatGPT." (reddit-103)
Firm-wide workflows; larger firm	Harvey or Legora over Claude Enterprise	The firm wanted packaged workflows and scheduling that ordinary users could adopt, not flexible one-off skills that still required configuration.	"The skills aren't real workflows, neither is the scheduling." (reddit-147)
Discovery, motion, and memo trial; about 30 attorneys	Internal build with Claude Opus over Harvey, Legora, and CoCounsel	Claude Opus found a material issue the specialist products missed. The firm also wanted its own DMS and project workflow and still needed Westlaw.	"The only one that caught [the heirship issue] ... was Claude Opus." (reddit-175)
Research authority plus generative drafting; size not stated	Claude plus CoCounsel/Westlaw connector over either alone	The buyer used CoCounsel for legal-source access and Claude for stronger generation and customization.	"Claude with custom skills and legal connectors ... works quite well at \$200/mo." (reddit-154)
Daily legal reasoning and drafting; size not stated	GC AI over Claude in one direct-use comment	The user wanted legal defaults and daily convenience without building a custom general-model workflow. Evidence remains sparse.	Reported result: GC AI was called "excellent for daily use" and preferred for legal reasoning. (reddit-167)

Sources: [reddit-110](#) | [reddit-131](#) | [reddit-103](#) | [reddit-147](#) | [reddit-175](#) | [reddit-154](#) | [reddit-167](#)

What the contrast shows

ChatGPT and Claude win when the buyer values model quality, breadth, customization, and low cost. Harvey and Legora win when packaged workflow and broad adoption matter more than model control. CoCounsel is strongest as an authority and litigation layer. GC AI may offer useful legal defaults, but the evidence is not yet strong enough to show a durable advantage.

Price decides the sale when the value gap is not obvious

Customers do not describe price as a separate purchasing factor. They reject the combination of a high fee and an answer that is not better than ChatGPT or Claude, a product that does not replace the rest of the stack, or a minimum commitment before value is proved. They accept a premium when the product removes measurable work. The next page lists every numeric price signal; this page shows the decisions in customers' own words.

Buyer and reported cost	What the customer said	Why the price won or lost
Large sub-AmLaw firm Harvey: \$1,200/seat/month; \$2,400 with Lexis; later about \$399	"At four figures a seat, I expect the demo folks to be able to answer basic questions with depth." It was not consistent. (reddit-117)	Harvey lost pricing power when the product demonstration did not show four-figure-per-seat performance. The same buyer reported about \$1,600/seat/month for its CoCounsel stack and said rewrite savings were not consistent.
Small company; exact size not stated Harvey: 5-seat minimum at \$425/seat/month for 12 months - a \$25,500 floor	"For 5 Harvey seats (minimum), 12 months lock-in, I was told it was USD \$425/seat/month." (reddit-163)	The issue was not only the seat price. A small buyer had to commit to five seats and a full year before proving use. Claude was the lower-cost comparison in the same discussion.
Full-service firm; about 100+ lawyers Harvey price not stated; Claude Enterprise plus a smaller research tool selected	"We actually tried Harvey early last year but its usefulness didn't justify the cost back then." (reddit-168)	Even an enterprise-size buyer rejected Harvey when usefulness did not clear the premium. Claude won on flexibility; the firm kept a separate legal-research layer.
Small California firm Harvey price not stated	"I was honestly shocked by how expensive and inflexible their pricing model is... no real path to try it." (reddit-107)	Inflexible terms and no meaningful small-team trial stopped the buyer before Harvey could prove workflow value.
Small firm; exact size not stated CoCounsel price increased; useful upgrades required more payment	"I use [CoCounsel] a lot... I'm considering switching to Westlaw 'Classic' with NO AI." (reddit-109)	A frequent user was willing to unbundle research and AI: keep lower-cost traditional Westlaw access, then buy a cheaper AI tool for generation.
Team size not stated Configured Claude with legal connectors: about \$200/month	"Claude with custom skills and legal connectors set up plus proper context works quite well at \$200/mo." (reddit-154)	Claude was 'good enough' because the buyer could add context and connectors while using CoCounsel mainly as a source layer, not as the generator.
Non-US firm; about 5,000 people Harvey rack-rate estimate near \$1,000/seat/month; enterprise discount received	"Saved me a hour or so there." The user said a 20-30% time saving could be "massive." (reddit-116)	This is the premium case: repeated time savings, easy use, and Vault document extraction can justify enterprise spend, especially on fixed-fee or capped work.

Direct sources: [reddit-117](#) | [reddit-163](#) | [reddit-168](#) | [reddit-107](#) | [reddit-109](#) | [reddit-154](#) | [reddit-116](#)

The pricing message

The specialist premium fails when the buyer still gets similar model output, keeps the existing research and document stack, and adds review or rewrite work. ChatGPT and Claude become the default because they are cheaper and broad enough. A specialist can still win when it produces repeated, measurable time savings or makes a large team easier to support.

Reported prices and the customer response

The figures below are anonymous, self-reported examples. They may refer to different dates, plans, regions, integrations, negotiations, or bundles. They are useful as evidence of buyer reaction, not as a current vendor price list.

Product	Reported price signal	Customer response and meaning
ChatGPT	\$20 consumer access was cited in one recent discussion; another small-firm user asked whether a \$200 monthly tier was worthwhile.	The low entry point makes ChatGPT the comparison baseline. Users accept more manual setup and verification because one subscription also supports nonlegal work.
Claude	One commenter described a configured Claude setup with legal connectors at about \$200 per month. Other posts called Claude a fraction of specialist cost without giving stable terms.	Customers see Claude as a way to buy model quality and flexibility first, then add only the research or workflow layer they need.
Harvey	Reports included about \$30,000 per year for a small legal department; a five-seat minimum at \$425 per seat per month for 12 months; and a large-firm quote of \$1,200 per user per month, \$2,400 with Lexis, later followed by a tier near \$399.	The spread creates pricing uncertainty. Small teams reject the premium when answer quality is not clearly better. Large firms can accept it when deployment, integration, sources, and document-scale work produce broad use.
Legora	The corpus contains no dependable numeric Legora price. Users often group it with Harvey as high enterprise spend, and one discussion objected to paying thousands per user for either product.	The price is difficult to defend when buyers see limited difference from Harvey or Claude. It is easier to defend when tabular review and standard workflows create repeated team savings.
CoCounsel	A 2023 report cited about \$50-\$75 per task and a subscription becoming preferable above roughly \$500 per month. A later frequent user objected to price increases and paid upgrades.	Task pricing can support a trial, but expanding fees weaken loyalty when users must still keep Westlaw and verify the output.
GC AI	The five GC AI records do not provide a usable numeric price benchmark.	Positive daily-use reports are not enough to determine whether buyers see good value versus ChatGPT or Claude.

Price evidence: [reddit-188](#) | [reddit-010](#) | [reddit-154](#) | [reddit-110](#) | [reddit-163](#) | [reddit-117](#) | [reddit-122](#) | [reddit-028](#) | [reddit-109](#)

What customers consider worth paying for

Worth a premium	Usually not worth a premium
Reliable large-file and batch review; clear review tables; precise redlines; legal-source access; DMS/Word links; approved data controls; firm-wide templates; low prompting burden; measurable repeated time savings.	A legal label alone; an answer similar to a general model; citations that still require full reconstruction; a chat interface without matter workflow; features used only occasionally; rigid output; extra seats; long lock-in before a real trial.

Cost conclusion

For small teams, ChatGPT or Claude is often the default because the downside is limited and the subscription has broad use. A specialist must prove that it removes enough verification, document handling, or workflow effort to exceed a large price gap. For large firms, the calculation changes: adoption, security, standardization, and integration can be worth more than the model itself.

ChatGPT: the broad, familiar baseline

ChatGPT is used as a general work tool rather than a complete legal system. Lawyers use it for first drafts, proofreading, brainstorming, issue spotting, clauses, correspondence, commercial advice, operations, and research framing. Its broad value makes the specialist premium difficult to justify for small teams.

What users like	What users dislike	Why it is chosen
Fast, familiar, broad across legal and nonlegal work, flexible instructions, strong writing and reasoning for some users, and low entry cost.	Fake cases, unreliable holdings, missed transcript or timeline details, weak final research, confidentiality concerns, and client-generated errors that create lawyer work.	Chosen when the team wants one versatile tool, can supply its own context, and will verify legal authority elsewhere.

Customer voices

Customer context	Reported voice or result	Source
GC; 3-attorney legal department	Preferred ChatGPT over Harvey and Claude on the same prompts for reasoning, citations, Word documents, and redlines; did not see about \$30,000/year of Harvey value.	reddit-110
Commercial-advice user; size not stated	Reported that ChatGPT Enterprise completed roughly 80% of a 25-page offshore-power advice task; review took about 10 hours instead of about a week of drafting.	reddit-167
Solo/small-firm lawyer	Used it for proofreading, brainstorming, and basic clause drafts after attorney editing; did not use it as a source of legal precedent.	reddit-010
Law-firm user; size not stated	Could not get ChatGPT to summarize an examination-before-trial transcript accurately.	reddit-009
SME legal department	Found enterprise controls useful for simple drafting and spotting missing items, but called legal research and substantive legal problem-solving failures.	reddit-031

Sources: [reddit-110](#) | [reddit-167](#) | [reddit-010](#) | [reddit-009](#) | [reddit-031](#)

Why users choose it over legal AI

ChatGPT often wins when the work is broad, the user is skilled, and authoritative research remains separate. It loses when confidentiality, reliable sourcing, document coverage, or firm-wide controls matter more than flexibility.

Claude: favored for drafting and configurable workflows

Claude is commonly used for transactional drafting, litigation workflows, contract review with playbooks, document synthesis, internal agents, coding, and repeatable projects. Recent customer voices often describe Claude as the strongest general-model alternative to a legal suite.

What users like	What users dislike	Why it is chosen
Strong drafting, large-context document work, projects and reusable instructions, custom skills and connectors, model choice, and good results when precedent or playbooks are supplied.	Usage limits, clunky Word-file generation, setup effort, no authoritative legal corpus, verification needs, and consumer-account privacy concerns.	Chosen by capable users and technology teams that want flexibility, lower cost, and control over models, precedent, and workflow.

Customer voices

Customer context	Reported voice or result	Source
Junior transactional lawyer; size not stated	Reported hours saved each week on defined terms, precedent cleanup, cross-references, and first drafts.	reddit-131
Partner commenter; size not stated	Said Claude Business plus agents outperformed a paralegal-and-associate team in a mock \$100 million loan workflow and was much better than CoCounsel in that test.	reddit-131
Law-firm technology lead; about 30 attorneys	Claude Opus caught a material heirship issue that Claude Sonnet, Harvey, Legora, and CoCounsel missed.	reddit-175
Technology committee; 100+ lawyer firm	Selected Claude Enterprise plus a smaller research tool after an earlier Harvey trial did not justify cost.	reddit-168
Plaintiff-employment lawyer	Reported a surprisingly decent case/discovery result but reached usage limits quickly.	reddit-042
In-house contract team	Separate Claude chats with a contract, playbook, and context produced useful work, but the team wanted a repeatable process instead of manual upload and email handoff.	reddit-132

Sources: [reddit-131](#) | [reddit-175](#) | [reddit-168](#) | [reddit-042](#) | [reddit-132](#)

Why users choose it over a legal suite

Claude wins when the buyer believes the model and custom context do most of the useful work. It loses when the team needs turnkey legal sources, review tables, Word/DMS integration, or low-skill firm-wide adoption.

Harvey: enterprise workflow value with contested small-team economics

Harvey has the broadest specialist discussion in the corpus. Positive users describe legal sources, review tables, large-document work, Word and DMS integration, Vault workflows, and firm-wide deployment. Negative users question whether those benefits justify the price when ChatGPT or Claude performs the underlying thinking and drafting well enough.

What users like	What users dislike	Why it is chosen
Turnkey legal workflows, easier prompting, review tables, document scale, redlining, legal-source links, Word/DMS integration, enterprise deployment, and shared standardization.	High or unclear pricing, minimum seats or commitments, mixed task quality, incomplete issue spotting, wrapper perception, and poor fit for small or uneven teams.	Chosen by larger organizations that value deployment, integration, and adoption across many lawyers more than direct model control.

Customer voices

Customer context	Reported voice or result	Source
Lawyer; about 5,000-person firm	Used Harvey regularly, reported about one hour saved on a pitch and good property-diligence results, but called it too expensive for small firms.	reddit-116
Mid-level BigLaw lawyer	Called Harvey life-changing for fast regulatory and clearance work; other commenters in the same thread reported limited use.	reddit-087
Corporate legal department; size not stated	Selected Harvey for nuanced examples, references, interface, Word integration, and legal databases despite no consistent answer-quality lead over ChatGPT.	reddit-103
GC; 3-attorney legal department	Rejected the value at roughly \$30,000/year after preferring ChatGPT on the tested work.	reddit-110
Small California firm	Rejected expensive, inflexible pricing and the lack of a meaningful small-team trial or customization path.	reddit-107
Law-firm technology lead; about 30 attorneys	Rejected Harvey with Legora and CoCounsel after all three missed a material issue and created DMS/project friction.	reddit-175

Sources: [reddit-116](#) | [reddit-087](#) | [reddit-103](#) | [reddit-110](#) | [reddit-107](#) | [reddit-175](#)

Why it wins or loses

Harvey wins when operational deployment is the product. It loses when a capable small team sees the legal workflow as an expensive layer around a general model it can already use.

Legora: strong review and drafting reports, but limited differentiation

Legora is discussed mainly in BigLaw and legaltech threads about research, drafting, diligence, tabular review, templates, and Word workflow. Positive users report frequent use and meaningful efficiency. Negative users see similar answers to Harvey or weaker output than enterprise Claude or ChatGPT.

What users like	What users dislike	Why it is chosen
Tabular review, diligence, drafting, research starting points, templates, Word workflow, standardized team use, and lower prompting burden.	High price, unclear difference from Harvey, mixed model quality, rigid output, missed issues, and concern that general models are catching up quickly.	Chosen when a firm wants a repeatable transactional or review workflow that many lawyers can use consistently.

Customer voices

Customer context	Reported voice or result	Source
BigLaw research discussion	One user said recent Legora use had been highly helpful for research; others still preferred constrained databases or traditional search because narrative explanations can hallucinate.	reddit-084
Large-firm discussion	One commenter reported using Legora every day for almost every legal task, while another mainly used Westlaw AI despite access to enterprise models.	reddit-088
BigLaw lawyer	Found Harvey and Legora similar on one research question and slightly preferred Harvey's clearer review-table categories; commenters also reported strong Legora drafting and a claimed 30% efficiency gain.	reddit-097
Law-firm technology lead; about 30 attorneys	Legora missed the same material issue as Harvey and CoCounsel in the three-week comparative trial.	reddit-175
Technology committee; 100+ lawyer firm	Had not trialed Legora and expected limited differentiation; this is expectation, not direct product evidence.	reddit-168
UK lawyer discussion	A user described extensive Legora tabular review for document questions such as identifying which document says a fact.	reddit-202

Sources: [reddit-084](#) | [reddit-088](#) | [reddit-097](#) | [reddit-175](#) | [reddit-168](#) | [reddit-202](#)

Why it wins or loses

Legora wins where review tables and standardized legal workflow produce visible team efficiency. It loses when the buyer cannot see a durable advantage over Harvey or a configured general model.

CoCounsel: strongest around authority and litigation workflow

CoCounsel is most often discussed as a Westlaw-connected research and litigation tool. Users apply it to discovery, depositions, records, timelines, brief review, and source-linked research. Its position is less secure for broad drafting and generative work.

What users like	What users dislike	Why it is chosen
Westlaw connection, source-linked starting points, legal safeguards, large uploads, discovery and deposition workflows, records review, and known legal-research deployment.	Generative limits, poor task following, missed concepts, timeline noise, weak full-document drafting, price increases, and continued citation checking.	Chosen when research authority, litigation workflow, or approved legal data matters more than open-ended drafting flexibility.

Customer voices

Customer context	Reported voice or result	Source
Firm user; size not stated	Praised legal-specific safeguards and pay-per-use options, but the firm paused a deposition use case over HIPAA certification concerns.	reddit-028
Small-firm lawyer	Reported using CoCounsel frequently but considered switching after price increases and paid upgrades weakened the value.	reddit-109
Paralegal and support users	Called discovery requests acceptable but reported that an amended complaint failed; records review missed requested concepts and created noisy timelines.	reddit-187
Westlaw research user	Preferred Westlaw Deep Research to the bundled CoCounsel experience and still found it useful even when it was seriously wrong; treated it as a checked starting point.	reddit-047
Claude/connector commenter	Said Westlaw plus CoCounsel worked as a connector with Claude, but CoCounsel was weak for generative work.	reddit-154
Law-firm technology lead; about 30 attorneys	Rejected CoCounsel with Harvey and Legora after the products missed a material issue and did not replace Westlaw.	reddit-175

Sources: [reddit-028](#) | [reddit-109](#) | [reddit-187](#) | [reddit-047](#) | [reddit-154](#) | [reddit-175](#)

Why it wins or loses

CoCounsel wins through the Westlaw relationship and litigation-specific work. It loses when users want stronger generative drafting, exact instruction following, lower cost, or a flexible model-centered workflow.

GC AI: promising daily legal use, but the evidence is sparse

GC AI appears in only five records. The available customer voices are positive about legal reasoning and drafting, but the sample is too small to establish a stable product position, target segment, or comparative advantage.

What users like	What users dislike	Why it is chosen
Legal-specific defaults, daily-use convenience, drafting, legal reasoning, playbooks, contract context, and repository-oriented workflow are reported positively by some users.	Too few independent reports, unclear firm-size fit, limited direct trials, polarized views, and uncertainty about whether the paid layer beats Claude or ChatGPT with good context.	Chosen by users who want legal context and workflow without building their own general-model setup, based on the limited evidence available.

Customer voices

Customer context	Reported voice or result	Source
Legaltech commenter; size not stated	Called GC AI excellent for daily use and preferred it to Claude for legal reasoning.	reddit-167
UK lawyer discussion; varied sizes	A participant reported using GC AI for drafting while others used Claude, Legora, Copilot, or bespoke tools for different jobs.	reddit-202
Frequent legal-AI evaluator; size not stated	Said most tested tools were not worth it but placed GC AI among the few products worth considering.	reddit-104
Broader market discussion	Other users argued that general models with web search, skills, and proper context handle much everyday legal work, making the specialist premium hard to prove.	reddit-167

Sources: [reddit-167](#) | [reddit-202](#) | [reddit-104](#)

Evidence warning

The available evidence supports a hypothesis that GC AI can be a useful legal daily driver. It does not support a confident conclusion about product-market fit, price-to-value, or superiority over Claude or ChatGPT.

What the six products reveal together

- General models own breadth, flexibility, and price-to-value.
- Specialist tools own selected workflows, controlled sources, integrations, and deployment.
- No product has established a consistent lead in legal reasoning across the direct comparisons.
- The buyer's team size, technical ability, existing stack, and verification burden often matter more than the product brand.

Direct selection and rejection decisions

The table below concentrates the strongest buying evidence. It shows why a product that works well for one organization can be rejected by another.

Customer / size	Products considered	Decision or outcome	Reason
GC; 3-attorney legal department	Harvey, ChatGPT, Claude	Preferred ChatGPT; rejected Harvey value.	Better tested reasoning, citations, Word and redline work; no reason to pay about \$30,000/year.
Technology lead; about 30-attorney firm	Harvey, Legora, CoCounsel, Claude Opus	Rejected all three specialists; chose internal build.	Material issue missed; DMS and project friction; products did not replace Westlaw.
Technology committee; 100+ lawyer firm	Claude Enterprise, Harvey, local research tool	Selected Claude Enterprise plus smaller research tool.	Earlier Harvey trial did not justify cost; Claude offered flexibility.
Lawyer; about 5,000-person firm	Harvey and alternatives	Used Harvey regularly and reported meaningful value.	Easy use, time saved, document-scale diligence, and enterprise availability.
Corporate legal department; size not stated	Harvey, ChatGPT	Selected Harvey despite no consistent answer lead.	References, examples, interface, Word integration, databases, and deployment.
Small California firm	Harvey	Rejected Harvey.	Price, inflexibility, no small-team trial, and weak customization path.
Small firm; size not stated	CoCounsel and alternatives	Considered leaving after frequent use.	Price increases and paid upgrades reduced value.
BigLaw discussion	Harvey, Legora, Claude, ChatGPT, Copilot	Mixed; no product winner.	Review-table clarity, drafting quality, research controls, and user preference differed by task.

Primary links: [reddit-110](#) | [reddit-175](#) | [reddit-168](#) | [reddit-116](#) | [reddit-103](#) | [reddit-107](#) | [reddit-109](#) | [reddit-097](#)

Firm size changes the value equation

Small teams tend to compare a specialist product with one or two inexpensive general-model seats and an existing research subscription. Large firms compare it with the cost of inconsistent work across hundreds of users, security review, training, DMS access, and firm-wide standardization. This is why high specialist pricing can be rational in a large deployment and unacceptable in a small department.

Important limit

Only about 22 corpus records identify a numeric team size or a solo practice. The size pattern is strong enough to guide interviews, but not to estimate demand or market share.

Where the current market still falls short

Market gap	Customer problem	Why current options do not fully solve it
Fast verification	Lawyers must check citations, holdings, facts, and source support before relying on the answer.	General models can invent sources; specialists can still be wrong or incomplete; source links do not prove the proposition was reviewed.
Proved document coverage	Users do not know whether every file, page, concept, or exception was reviewed.	Long context windows and large-upload claims do not produce a clear coverage receipt or exception list.
Reliable Word handoff	Copy-paste, broken formatting, weak redlines, and unexplained edits create a second review job.	General models are clunky; specialist Word tools vary; source-layout preservation remains difficult.
Persistent context	Facts, prior decisions, playbooks, open questions, and work product are rebuilt across chats and tools.	Projects help, but users still report manual setup, separate chats, and disconnected systems.
Small-team economics	Solo and small teams need value without enterprise minimums, rollout, or unused seats.	Many specialist products are designed and priced for larger firms; point tools fragment the workflow.
Flexible but usable workflow	Experts want customization; ordinary users need a turnkey workflow.	General models require skill and setup; specialist products can be rigid or hide model choice.
Authority plus drafting	Users want controlled research and strong generative drafting in one flow.	Research tools are weaker generators; general models lack controlled authority; users bridge them manually.
Matter follow-through	AI produces an answer but does not reliably preserve the decision, owner, deadline, and next action.	Chat and point tools often stop at output rather than moving the legal matter forward.

The central market tension

Customers want the intelligence and flexibility of the newest general models, but they also want the sources, document control, integrations, and repeatability of a legal workflow product. The market has not settled on one product that delivers both without high cost, technical setup, or continued manual verification.

Unmet job

Turn a strong model answer into dependable legal work: source-linked, complete enough to review, connected to the actual documents, and preserved with the matter's next action.

Comparative market summary

Product	Strongest observed position	Main reason to choose	Main objection
ChatGPT	Broad general work	Versatility, familiarity, price, and useful drafting/reasoning.	Research reliability, confidentiality, and no durable legal workflow.
Claude	Drafting and configurable document workflow	Writing quality, projects, skills, context, and model control.	No authoritative corpus, usage/setup friction, and weak Word handoff.
Harvey	Enterprise legal workflow	Review tables, document scale, sources, integration, and adoption.	Price, commitments, and mixed advantage over general models.
Legora	Review, diligence, and standardized drafting	Tabular review, Word workflow, and team consistency.	Unclear differentiation, price, and mixed model-quality reports.
CoCounsel	Research and litigation	Westlaw, legal sources, discovery, records, and deposition workflow.	Generative limits, task failures, price, and continued checking.
GC AI	Legal daily-use assistant	Legal reasoning, drafting, playbooks, and convenience for some users.	Too little evidence to establish a clear moat or buyer segment.

Short comparison: Themis / CounselOS

Themis is currently closest to the persistent-matter and follow-through gap. Its strengths are model neutrality, matter context, editable work product, decisions, work items, and visible next action. It is not a replacement for Harvey or Legora's enterprise deployment, CoCounsel's controlled legal research, or a live Word/DMS workflow.

Current overlap with market needs	Material gaps
Durable matter workspace; model choice; company and matter context; editable Markdown work product; research packets; decisions; work items; schedules; DOCX/PDF export.	Exact passage-level proof; document-coverage receipts; structured batch review; source-layout-preserving or live Word workflow; authoritative Westlaw/Lexis handoff; team access; email/Slack/DMS integration; enterprise controls.

Market-research conclusion

Themis addresses a real gap identified in the customer voices, but the evidence does not yet show that customers will pay for its approach. The next proof should be whether it reduces repeated setup and verification on real matters better than a user's existing ChatGPT or Claude stack.

Evidence strength and limitations

The corpus contains 203 de-duplicated Reddit threads. The latest expansion added 100 unique threads dated March through August 2026. Direct product use, comparative trials, reported selection decisions, observable work, and stated team size receive the most weight. Broad opinions, vendor claims, affiliate comments, hearsay prices, and removed posts receive less weight.

- A product mention is not a positive review. Counts show discussion coverage only.
- A Reddit report is not an independent product benchmark. Model versions, product releases, prompts, files, and user skill differ.
- Self-reported pricing may reflect negotiation, region, plan tier, database bundles, or incomplete terms.
- Firm and legal-team size is usually missing. Size conclusions are directional.
- The report describes customer voices and market gaps. It does not estimate market size or market share.

Curated source appendix

The links below are the primary sources used for the product profiles and cross-product conclusions. The complete machine-readable corpus remains at [output/research/raw/legal-ai-reddit-source-corpus.jsonl](#).

ChatGPT and Claude

- [PERIODIC CHECK-IN: Has anyone figured out how to actually use ChatGPT or any other LLM to actually make their jobs easier / better / faster? \(reddit-009\)](#)
- [ChatGPT? \(reddit-010\)](#)
- [ChatGPT Enterprise / Inhouse Lawyer - Thoughts and Discussion \(reddit-031\)](#)
- [Have any in-house counsel found practical and accurate uses of AI for day-to-day legal work? \(reddit-035\)](#)
- [Claude Cowork \(reddit-042\)](#)
- [Claude Litigation workflows? \(reddit-057\)](#)
- [Harvey for In-house Counsel \(reddit-110\)](#)
- [Using Claude for drafting transactional documents \(reddit-131\)](#)
- [How to build a contract review process using Claude for an in-house team? \(reddit-132\)](#)
- [What's your pricing on Harvey or Legora? \(reddit-154\)](#)
- [Reasons to reject Harvey / Legora \(reddit-168\)](#)
- [We trialed CoCounsel, Harvey, and Legora. Here is why we're building our own internal tool instead. \(reddit-175\)](#)

Harvey and Legora

- [Does anyone like AI for legal research? \(reddit-084\)](#)
- [Harvey \(reddit-087\)](#)
- [Am I the only one who isn't seeing widespread AI adoption and has no idea what the posts here are on about? \(reddit-088\)](#)
- [Harvey vs Legora \(reddit-097\)](#)
- [Harvey \(reddit-103\)](#)
- [Harvey AI Says It's for All Lawyers — But Prices Like It's Only for Biglaw \(reddit-107\)](#)
- [Harvey AI \(reddit-116\)](#)
- [Pricing: Harvey v. Claude v. Legora v. CoCounsel \(from what we were quoted\) \(reddit-117\)](#)
- [Law firms: How much do you pay for Legora/Harvey/similar? \(reddit-124\)](#)
- [Help me understand what Harvey has that Claude + connectors do not \(reddit-163\)](#)
- [Reasons to reject Harvey / Legora \(reddit-168\)](#)
- [We trialed CoCounsel, Harvey, and Legora. Here is why we're building our own internal tool instead. \(reddit-175\)](#)
- [Anyone finding AI useful? How are you using it? \(reddit-202\)](#)

CoCounsel and GC AI

- [AI in your practice \(reddit-028\)](#)
- [Westlaw AI Deep Research Users? \(reddit-047\)](#)
- [This Sub is Being Turned Into One Big Market Survey for Legal AI \(reddit-104\)](#)
- [Another Punch in the Gut From Westlaw/CoCounsel \(reddit-109\)](#)
- [What's your pricing on Harvey or Legora? \(reddit-154\)](#)
- [Are you seeing much AI at work? \(reddit-165\)](#)
- [Unpopular opinion: most Legal AI tools are not worth paying for. \(reddit-167\)](#)
- [We trialed CoCounsel, Harvey, and Legora. Here is why we're building our own internal tool instead. \(reddit-175\)](#)
- [Westlaw CoCounsel \(reddit-187\)](#)
- [Anyone finding AI useful? How are you using it? \(reddit-202\)](#)

Small-team and market context

- [What AI tools do you use \(not a bot!\) \(reddit-017\)](#)
- [How are you getting past the confidentiality fears with AI? \(reddit-025\)](#)
- [SmallLaw Lawyers, how do you use AI? \(reddit-050\)](#)
- [Tech stack for solos in 2026 \(reddit-054\)](#)
- [Client looking for advice \(reddit-119\)](#)
- [62 legal AI assessments in, here's what surprised me about small firm buyers \(reddit-129\)](#)
- [Are there Legal AI platforms for smaller transactional firms like Harvey/Legora? \(reddit-152\)](#)
- [How is everyone using AI to scale? \(reddit-195\)](#)